

How do the client's tactical instructions reach the proposal?

1. Pass through — verbatim — no sorting, no classifying

The client's tactical text goes straight into the prompt, unedited — together with the intake parameters, the parsed holdings + statement source, and the research / other documents in full.

2. One self-contained prompt — the transparency surface


Shown on the Proposal page — editable, copyable, downloadable. Paste it into ANY AI model to reproduce the proposal and compare outputs. Generate in-app runs the live AI model.

3. Numbers stay deterministic — the guardrail

Every figure is computed by the engine (the FACTS block). The AI writes prose only; qualitative claims from the documents are context, not independently verified.

What changed in v10: the copilot no longer sorts instructions into typed items, tags enforcement tiers, or gates a trade on a price trigger. It hands the raw inputs to the AI Model in one prompt and lets it analyse them — while the engine still computes every number, so nothing is invented. The UI is model-agnostic: it always reads "AI Model".

Tactical instructions — who does what, and when



Recommended order: set the policy (1–2) first, then Analyse (4). The tactical text and documents pass verbatim into one self-contained prompt (5–6) that the AI Model writes from — and that you can copy into any AI model. Every figure is computed deterministically: the AI writes prose only.